

In [10]: *## importing the required libraries*

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error
from sklearn.impute import SimpleImputer
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
```

In [11]: *## Load the data*

```
train_data = pd.read_csv('train.csv')

## Drop rows with missing target values
train_data.dropna(axis=0, subset=['SalePrice'], inplace=True)

## Define target (y) and features (X)
y = train_data['SalePrice']
X = train_data.drop(['SalePrice'], axis=1).select_dtypes(exclude=['object'])

## Now, ensure y is aligned with X after dropping missing values
y = y[X.index]
```

In [12]: *## Split the data into training and testing sets*

```
## 25% of the data will be used for testing, the rest for training
## Feature (x) and target (y) splits
train_X, test_X, train_y, test_y = train_test_split(X.values, y.values, test_size=0.25, random_state=42)

## Initialise the imputer (you can choose strategy: 'mean', 'median', or 'most_frequent')
my_imputer = SimpleImputer() # Default is 'mean', but can choose others

## Fit and transform the training data (impute missing values in the training set)
train_X = my_imputer.fit_transform(train_X)

## Transform the test data using the same imputer (to avoid data leakage)
test_X = my_imputer.transform(test_X)
```

```
## Optionally, you can convert back to DataFrame with the original column names after imputation  
train_X = pd.DataFrame(train_X, columns=X.columns)  
test_X = pd.DataFrame(test_X, columns=X.columns)
```

In [13]: *## Predictions using the Decision Tree algorithm*

```
decision_model = DecisionTreeRegressor()  
decision_model.fit(train_X, train_y)  
predicted_decision_trees = decision_model.predict(test_X)  
print("Mean Absolute Error using Decision Tress :", mean_absolute_error(test_y, predicted_decision_trees))
```

Mean Absolute Error using Decision Tress : 24394.780821917808

In [14]: *## Predictions using the Random Forest algorithm*

```
forest_model = RandomForestRegressor(n_estimators=100, max_depth=10)  
forest_model.fit(train_X, train_y)  
predicted_random_forest = forest_model.predict(test_X)  
print("Mean Absolute Error using Random Forest:", mean_absolute_error(test_y, predicted_random_forest))
```

Mean Absolute Error using Random Forest: 17481.22066004496

In [15]: *## Predictions using the XGBoost algorithm*

```
xg_model = XGBRegressor(n_estimators=100)  
xg_model.fit(train_X, train_y)  
predicted_XGBoost = xg_model.predict(test_X)  
print("Mean Absolute Error using XGBoost: ", mean_absolute_error(test_y, predicted_XGBoost))
```

Mean Absolute Error using XGBoost: 18624.489811643834

In [16]: *## Let's improve the RandomForest Model as it was the best performing in MAE*  
*## We will use GridSearchCV to tune the hyperparameters of the RandomForest Model*

```
from sklearn.model_selection import GridSearchCV  
  
## Define the hyperparameter grid to search  
param_grid = {  
    'n_estimators': [100, 200, 300],  
    'max_depth': [10, 20, 30, None],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 4],
```

```

    'max_features': ['sqrt', 'log2']
}

## Create a RandomForestRegressor model
rf_model = RandomForestRegressor()

## Perform grid search with cross-validation
grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid, cv=2, n_jobs=-1, verbose=2)
grid_search.fit(train_X, train_y)

## Get the best parameters and best score
print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

## Predict with the best model
best_rf_model = grid_search.best_estimator_
predicted_random_forest = best_rf_model.predict(test_X)
print("Mean Absolute Error using optimised Random Forest:", mean_absolute_error(test_y, predicted_random_forest))

```

Fitting 2 folds for each of 216 candidates, totalling 432 fits

c:\Users\samuel.mcdonnell\AppData\Local\anaconda3\envs\HousePrices\Lib\site-packages\numpy\ma\core.py:2820: RuntimeWarning: invalid value encountered in cast

```
_data = np.array(data, dtype=dtype, copy=copy,
```

```
Best parameters: {'max_depth': None, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 300}
```

```
Best score: 0.8350334919908053
```

```
Mean Absolute Error using optimised Random Forest: 16769.477760700804
```

```

In [17]: ## Get the feature importance scores
feature_importances = forest_model.feature_importances_

## Create a List of features with their corresponding importance scores
feature_importance_dict = {
    feature: importance for feature, importance in zip(X.columns, feature_importances)
}

## Sort the features by their importance score (in descending order)
sorted_features = sorted(feature_importance_dict.items(), key=lambda x: x[1], reverse=True)

## Display the top 10 most important features

```

```
top_10_features = sorted_features[:10]
print("Top 10 Features by Importance:")
for feature, importance in top_10_features:
    print(f"{feature}: {importance}")

## Random Forest inherently handles feature importance so it isn't necessary to include in tuning process. Still interesting t
```

Top 10 Features by Importance:  
OverallQual: 0.5624306342744222  
GrLivArea: 0.12098009777652968  
TotalBsmrSF: 0.03809611172661414  
2ndFlrSF: 0.03260436035765927  
1stFlrSF: 0.030584864830592924  
BsmrFinSF1: 0.029765704475803872  
GarageCars: 0.0265242748256848  
GarageArea: 0.022851367791071046  
LotArea: 0.020380813784723593  
YearBuilt: 0.015883262828223846

```
In [19]: mean_absolute_error(test_y, predicted_random_forest)/train_data['SalePrice'].median()
```

```
Out[19]: 0.10288023166074113
```

In conclusion, this Random Forest model, after hyperparameter tuning, achieved a Mean Absolute Error of 16,769. The median of the house price data is 163,000. This gives us a relative error of 10.3%, a great result considering the complexity of the data.